

Per-Word Arabic *I'rāb* Generation: A Comparison Study with a Cross-Register Audit

Hatem Saadallah

Nour Jennane

May 2026

Abstract

Arabic *i'rāb* (إعراب) generation produces structured grammatical analysis as natural-language prose. We report two findings that we believe should reframe the next research step on the task. First, **scale within Arabic-pretrained open-weight fine-tuning is teacher-bound**: three Arabic-pretrained models from 296M (AraT5v2-base) to 13B parameters (AceGPT-13B), trained on the same 77K-row Claude-Haiku-distilled corpus, reach statistically indistinguishable Gazelle performance — every pairwise McNemar $p=1.000$ at $n=134$ word judgments. The students inherit the teacher's case-67% ceiling; on this benchmark, additional capacity has no signal to exploit. Second, **register transfer breaks the strongest system the most**. Moving from MSA news (Gazelle) to Quranic (a 5,007-word MASAQ subset), every LLM-trained system loses paired-significant role-F1 (≥ 14 pp); the closed-source frontier — Claude Sonnet 4.5 + RAG, our Gazelle headline at 32.1% fully-correct [24.6, 40.3] — shows the *largest* cross-register drop of any system tested (+61.7 pp★). A retrieval-pool-confound experiment rules out the obvious explanation: removing the MSA retrieval pool (Sonnet zero-shot) widens the drop to +67.0 pp★ rather than narrowing it. Within the resolution this benchmark allows, register and construction coverage — not parametric capacity — account for the remaining error. We argue the cheapest informative next step is a Sonnet-distilled re-training (which would test whether the capacity null is teacher-bound) and targeted annotation of the constructions where every system silently fails.

1 Introduction

The *i'rāb* task takes an undiacritized Arabic sentence and produces per-word traditional grammatical analysis as Arabic prose. Each word receives four fields: part of speech, case (*raf'*/*naṣb*/*jarr*/*jazm*/*mabni*), syntactic role (~ 25 traditional labels), and marker phrase (~ 200 unique surface forms, e.g. على الظاهرة الضمة (آخره)). The task matters for two practical reasons. First, every Arabic-as-a-second-language curriculum uses *i'rāb* as the canonical pedagogy framing for Arabic syntax. An automatic *i'rāb* tool

would directly help language learning, religious-studies tooling, and corpus annotation. Second, the task probes fine-grained Arabic grammatical knowledge that POS accuracy or dependency UAS misses. The role + case + marker triple requires resolving subject/object selection, *idāfa* state, and exception (*istiṭhāna*) constructions whose interactions are non-trivial.

The pain. The literature on Arabic *i'rāb* reads as two non-comparable threads. Classifier pipelines (CamelParser2.0, MADAMIRA, Farasa) report dependency / morphological accuracy on tag-level outputs. Frontier-LLM evaluations (e.g. Gazelle) report multiple-choice or string-similarity scores on prose. Two questions remain unanswered. (i) When the same system is asked to *generate* prose *i'rāb*, do extra parameters help, or does the training data ceiling dominate? (ii) Does a system that scores well on MSA news still score well on a different register (e.g. Quranic), or does the headline number hide a register-specific competence?

The gain. We answer both questions on a controlled testbed. A regex/FSM extractor parses every prediction into {case, role, marker}, and every system-vs-system delta carries a paired bootstrap CI and a McNemar exact p -value (the methodology is infrastructure for the findings, not a contribution in itself). Within this framework we evaluate eleven systems — one classifier-pipeline baseline (Stanza), one open-weight LLM peer (Qwen2.5-7B+RAG), four open-weight fine-tuned models from 296M to 13B trained on the same Haiku-distilled corpus, four closed-source LLM variants (Haiku and Sonnet, zero-shot and +RAG), and one routing hybrid (Mix A). We then re-evaluate the six trainable / closed systems on a held-out 5,007-word Quranic MASAQ subset.

Contributions.

1. **Teacher-bound scaling on *i'rāb* fine-tuning** (§3.2, Fig. 1). Three Arabic-pretrained open-weight models spanning 296M to 13B parameters, trained on a single Claude-Haiku-distilled corpus, reach statistically




Figure 1: Gazelle *fully-correct* rate vs. open-weight parameter count (log scale). Three Arabic-pretrained models (296M, 792M, 13B; blue) train to statistically indistinguishable performance — every pairwise McNemar $p=1.000$. The closed-system frontier (Sonnet RAG, dashed red) sits 7–8 pp above. Multilingual mT5 (grey) trails by paired-significant margin. Scale within Arabic-pretrained open weights does not move the headline.

indistinguishable Gazelle performance (every pairwise Δ has McNemar $p=1.000$). The teacher’s $\sim 67\%$ case ceiling is the binding constraint, not the student’s capacity. We treat this as a constraint specific to Haiku-as-teacher and the Gazelle benchmark size, and we propose a Sonnet-redistillation experiment that would directly test whether scale begins to matter once the teacher ceiling lifts.

2. **Cross-register fragility, peaked at the strongest system** (§3.5). The closed-source frontier on MSA news (Sonnet RAG, +61.7 pp★ drop, retains 19% of MSA role-F1 on Quranic) is the system most disturbed by the register shift, not the most robust. A confound experiment rules out the natural “the retrieval pool is MSA” explanation: zero-shot Sonnet drops *more*, not less. Within the resolution of this benchmark, the remaining error is dominated by register and construction coverage rather than by model capacity, which reframes the next research step toward targeted annotation rather than larger models.

2 Experiments

2.1 Data

We use three categories of *i’rāb* data, distinguished by their level of human authorship.

Manual gold (held-out evaluation).

Gazelle [5]: 30 hand-written MSA sentences ($n=134$ word judgments). *MASQA* [12]: 624 Quranic verses ($n=5,007$ word judgments); the 999-word subset where the gold contains an extractable role label provides a matched cross-register surface. *Yarob* (linuxscout/yarob, GPL):

459 MSA *i’rāb* sentences in classical-grammar prose; we use these as the retrieval pool, not as evaluation.

LLM-distilled training corpus. We use Claude Haiku 4.5 as a teacher to generate per-word *i’rāb* for 5,000 PADT MSA news sentences via the Anthropic Messages Batches API. After JSON-schema validation, we obtain **77,534 (word, sentence, irab) rows**. This is the training corpus for every open-weight FT model below; no model trains on Gazelle or MASQA.

Construction coverage. A regex over each Gazelle gold prose tags the dominant construction(s) per sentence using 11 traditional categories. The 30-sentence eval is dominated by short verbal sentences ($n=15$) and prepositional phrases ($n=8$); exception (*istithnā*) and *kāna* sisters appear in 2 sentences each ($n=9$ and $n=7$ words). Generalization claims are limited to these construction frequencies.

2.2 Models

We evaluate eleven systems on Gazelle.

Classifier-pipeline baseline. *Stanza Arabic* [10] with a deterministic UD→*i’rāb* stub (the published-prior-work peer).

Open-weight LLM peer. *Qwen2.5-7B-Instruct* + RAG ($k=5$, 4-bit local).

Open-weight FT models (same 77,534-row Haiku corpus, same prompt template):

- AraT5v2-base (296M, T5 enc-dec) [8] — full FT, 3 epochs.
- mT5-base (580M, T5 enc-dec) [13] — full FT, 3 epochs.
- AraGPT2-large (792M, GPT-2 decoder-only) [3] — LoRA $r=32$, 3 epochs.
- AceGPT-13B (13B, Llama-2 decoder-only) [6] — QLoRA 4-bit $r=32$, 1 epoch.

Closed-source LLMs. Claude Haiku 4.5 (zero-shot, +RAG $k=5$); Claude Sonnet 4.5 (zero-shot, +RAG $k=5$). RAG uses Jaccard retrieval over a 1,060-example pool (Yarob + a 601-row distilled prior) following the standard recipe [7].

Hybrid Mix A. The LLM-RAG output supplies case and role; the AraT5v2-base specialist supplies the marker phrase, called per word with the structured prompt `[case=c] [role=r] w | s → marker`.

A six-system subset (Stanza, mT5, AraT5v2, AraGPT2, AceGPT-13B, Sonnet RAG) also runs on MASQA; the AceGPT-13B MASQA run hit a 4 h SLURM timeout at 21% of MASQA ($n=1,075$ words), and we report it with that caveat.

2.3 Setup

Structural metric. A regex/FSM extractor parses each generated $i'rāb$ string into four fields: pos, case $\in \{raf^c, naṣb, jarr, jazm, mabni\}$, role (~ 25 labels), marker (~ 200 phrases). Macro role-F1 over a closed label set L is

$$\text{role-F1}_{\text{macro}} = \frac{1}{|L|} \sum_{\ell \in L} \frac{2P_{\ell}R_{\ell}}{P_{\ell} + R_{\ell}}, \quad (1)$$

where P_{ℓ}, R_{ℓ} are precision and recall for label ℓ . We additionally report **case-acc**, **marker-EM**, **well-formed**, and the headline aggregate

$$\text{fully}(w) = \mathbf{1}[\text{case}(w)] \cdot \mathbf{1}[\text{role}(w)] \cdot \mathbf{1}[\text{marker}(w)]. \quad (2)$$

We avoid chrF and BLEU because partial-credit on string overlap masks morphological errors that we explicitly target.

Statistical inference. Every per-system metric carries a 95% percentile bootstrap CI ($B=1000$):

$$\text{CI}_{0.95}(\theta) = [\theta_{(0.025B)}^*, \theta_{(0.975B)}^*]. \quad (3)$$

For paired comparisons of two systems on the same n word judgments, we report the paired-bootstrap CI of $\Delta = \theta_A - \theta_B$ and McNemar’s exact p -value on the discordant-pair count b/c :

$$p_{\text{McNemar}} = 2 \min \left(\sum_{k=0}^{\min(b,c)} \binom{b+c}{k} 2^{-(b+c)}, 0.5 \right), \quad (4)$$

computed in log-space for $n > 1000$. A Δ counts as significant ($\star$) only if its 95% paired CI excludes 0 and $p_{\text{McNemar}} < 0.05$. Cross-register comparisons (Gazelle subset $n=78$ vs. MASAQ subset $n=999$) use a two-sample bootstrap, since the surfaces share no items. The smallest detectable difference at $n=134$, $\alpha=0.05$, power 0.80 is approximately ± 7 pp on a binary metric.

Extractor audit. We perturb each Gazelle gold record in three ways (case-flip, role-flip, marker-mangle), producing 402 records, and re-score with the same extractor. Specificity reaches 100% on unperturbed controls (no false positives); sensitivity is 92.7% (marker), 88.2% (case), 60.0% (role). The 60% role floor reflects the closed taxonomy of ~ 25 traditional labels; we read role results with this caveat.

3 Results

3.1 Headline (Gazelle)

Table 1 shows the headline. Three findings follow.

3.2 Three null results from the comparison

(a) Open-weight scale is null on Gazelle. AraT5v2-base (296M), AraGPT2-large (792M), and AceGPT-13B (13B) — three architectures (T5 enc-dec, GPT-2 decoder-only, Llama-2 decoder-only), $44\times$ parameter range — train to statistically indistinguishable Gazelle performance. Every pairwise Δ has McNemar $p=1.000$ (e.g. AceGPT-13B vs. AraT5v2-base: case $\Delta +0.7$, marker $\Delta -0.7$, fully $\Delta +0.0$). Fig. 1 visualises the result.

(b) Mix A’s per-word routing does not significantly improve over Sonnet RAG. Δ case $+0.0$ pp ($p=1.000$), Δ marker -3.7 pp ($p=0.125$), Δ fully -3.0 pp ($p=0.219$). The routing hypothesis — that case/role decompose from marker phrasing — does not hold here. Two factors plausibly explain why. The marker specialist trains on 8,815 marker pairs of which 7,172 (81%) are themselves Haiku-distilled; the specialist therefore learns “what Haiku would have said” rather than an independent gold style, and RAG already exposes Sonnet to that same demonstration distribution. Additionally, the Sonnet-base hybrid rebuild template truncates longer LLM role labels (e.g. $\text{إلى الجرح بحرف مجرور اسم}$ collapses to مجرور اسم), which slightly hurts role-F1. The result also holds on the Haiku base (Δ fully -1.5 pp, $p=0.791$); we report Mix A as a documented negative result on the routing decomposition at the present evaluation scale.

(c) Arabic-specific pretraining at 296M beats multilingual pretraining at 580M. AraT5v2-base – mT5-base: Δ marker $+5.2$ pp $\star$ ($p=0.039$), Δ fully $+6.0$ pp $\star$ ($p=0.021$). Same training corpus, same recipe, same architecture — only the pre-training corpus differs.

The closed-system advantage is real but bounded: Sonnet RAG beats AraT5v2-base by Δ case $+8.2$ pp $\star$ and Δ fully $+7.5$ pp $\star$. Open-weight scale alone does not narrow this gap.

3.3 Per-construction failures

Fig. 2 surfaces two cross-system structural failures: EXCEPTION (0/9 across every evaluated system) and KANA SISTERS (0/7 across four of five). Both failures hold for the trained AraT5v2-base and AceGPT-13B as well as for every Claude variant. We read these as construction-frequency artefacts rather than as a Claude-specific stylistic gap. Exception ($isti^{th}nā$) requires recognising that the post-marker noun (after إلا/عدا/سوى) takes case based on whether the exception is positive or negative — a rule that appears rarely in news corpora and almost never in our 77K Haiku-distilled training rows. *Kāna* sisters require shifting the predicate to $naṣb$; the Haiku teacher itself

Table 1: Headline comparison on Gazelle ($n=134$ word judgments). All numbers are percentages; brackets are 95% percentile bootstrap CIs ($B=1000$) shown for the headline metric *fully* only. Boldface = paired-best within column. The from-scratch decoder appears in Appendix A.

System	well	case	role-F1	marker	fully
Stanza Arabic (UD+stub)	59.7	35.1	10.3	13.4	5.2 [2.2, 9.7]
Qwen2.5-7B-Instruct+RAG	66.4	43.3	19.2	19.4	3.0 [0.0, 6.0]
mT5-base (580M) FT	79.9	61.9	31.6	32.8	18.7 [11.9, 25.4]
AraT5v2-base (296M) FT	79.9	65.7	54.8	44.0	24.6 [17.9, 32.8]
AraGPT2-large (792M) LoRA FT	79.9	64.9	54.6	43.3	26.1 [19.4, 34.3]
AceGPT-13B QLoRA FT (1 epoch)	79.1	66.4	55.5	43.3	25.4 [17.9, 32.8]
Claude Haiku 4.5 zero-shot	77.6	57.5	55.9	40.3	18.7 [11.9, 25.4]
Claude Haiku 4.5 + RAG	79.9	67.2	68.9	44.8	27.6 [20.1, 35.8]
Mix A (Sonnet RAG + AraT5v2 marker)	79.9	73.9	73.3	46.3	29.1 [21.6, 37.3]
Claude Sonnet 4.5 zero-shot	78.4	72.4	76.4	44.0	27.6 [20.1, 35.8]
Claude Sonnet 4.5 + RAG	79.9	73.9	74.7	50.0	32.1 [24.6, 40.3]




Figure 2: Per-construction *fully*-correct rate on Gazelle. EXCEPTION ($n=9$) and KANA_SISTERS ($n=7$) collapse to 0% across every system, including the strongest. Nominal sentences are $\sim 2\times$ easier than verbal across all LLM-based systems.

does not reliably output منصوب خبر, so distilled students inherit the failure. The point estimates are robust to bootstrap noise (zero cannot widen) but *not* to coverage of these constructions in training data; we expect a few hundred targeted examples per construction to break the floor. The two sentence-level samples ($n=9$ and $n=7$ words) are too small to generalise to all instances of the construction; we report the failure as a clear signal at this scale, not as a population-level claim.

3.4 Audits and ablations

System discrimination. When Sonnet RAG verifies perturbed gold against the source, its case-acc collapses from 96% to 15% ($\Delta = -81$ pp). The metric reacts sharply to incorrect inputs. **Retrieval depth k .** Sonnet RAG case-acc varies $70.1 \rightarrow 73.9 \rightarrow 73.1$ across $k \in \{1, 3, 5, 8, 12\}$; *fully* varies $28.4 \rightarrow 32.1 \rightarrow 30.6$. No paired-significant difference between $k=3$ and $k=8$ ($p=0.625$ on *fully*); $k=5$ is not a free parameter doing meaningful work. **Inference-time variance ($T=0$).** Two repeats agree on 96.8% of words; case-acc differs by 0.7 pp. **Annotator disagreement.** A second annotator finds 31% of Gazelle words admit an alternative-but-defensible analysis; permissive scoring lifts Sonnet RAG case by only




Figure 3: Cross-register Δ subset role-F1 (Gazelle–MASAQ subset, in pp). The strongest Gazelle system (Sonnet RAG, red) shows the *largest* cross-register drop. Stanza (UD pipeline, grey) is the only register-stable system. Bars: 95% two-sample bootstrap CIs; $\star$ marks paired-significant deltas.

Table 2: Subset role-F1 (gold has extractable role; matched cross-register definition). 95% two-sample bootstrap CIs, $B=1000$.

System	Gazelle	MASAQ	Δ (pp)
Stanza	10.4	17.6	-7.2 (ns)
mT5-base	32.8	18.6	+14.2 $\star$
AraT5v2-base	58.9	24.3	+34.7 $\star$
AraGPT2-large	58.1	20.2	+37.9 $\star$
AceGPT-13B	60.4	22.2 $\star$	+38.2 $\star$
Sonnet RAG	75.7	14.1	+61.7 $\star$
Sonnet zero-shot**	78.1	11.0	+67.0 $\star$

*partial MASAQ ($n=210$); **first 400 verses ($n=657$).

+0.7 pp. The headline is robust to plausible label noise.

3.5 Cross-register audit

Table 2 and Fig. 3 report three findings. (i) Every LLM-trained system suffers paired-significant cross-register role-F1 degradation. (ii) Stanza (UD pipeline) is register-stable; the closed UD POS+deprel label set applies uniformly across registers. (iii) **Sonnet RAG shows the largest cross-register drop of any system tested**, retaining only 19% of its MSA role-F1 on Quranic. Scale ($44\times$) and architecture do not narrow this

gap.

Retrieval-pool confound: rejected. Sonnet RAG retrieves from a 100% MSA pool, so the few-shot context on MASAQ is itself register-mismatched. We re-run Sonnet zero-shot on the first 400 MASAQ verses ($n_{\text{subset}}=657$): zero-shot subset role-F1 = 11.0, *lower* than RAG’s 14.1 (CI $[-2.1, +6.8]$ crosses 0); the cross-register Δ for zero-shot is +67.0 pp $\star$, *larger* than RAG’s. Removing retrieval widens the drop, not narrows it. The MSA-pool retrieval is therefore not what is masking cross-register competence; the LLM’s parametric knowledge is itself register-specific.

Why the strongest system drops the most. A natural reading of Tab. 2 is that Sonnet RAG simply has more to lose: its MSA ceiling is highest, so the absolute gap to its Quranic floor is largest. The percentage figures push back on the pure “regression-to-mean” story — Sonnet RAG retains only 19% of its MSA role-F1, while AraT5v2-base retains 41% and Stanza retains 169% (it actually goes *up* on Quranic). Stanza’s stability is informative: its closed UD POS+*deprel* label set, mapped via the same templater on both surfaces, removes the lexicalised “what does this Arabic word usually do” signal that the LLMs lean on. The LLMs leverage MSA-news vocabulary distributions and constructions (verb-initial copular clauses, news-style *idāfa*); Quranic surface forms (rhymed verse, classical syntactic devices, lexical items not in modern news) sit outside this distribution, and a richer parametric prior (Sonnet) is more confidently wrong than a weaker one (Stanza, mT5). We do not claim this explanation is exclusive — a careful per-token error breakdown across 5,007 MASAQ words would test it directly — but it is consistent with every cross-register row in Tab. 2.

4 Related Work

Arabic morphosyntactic tools. Classifier-pipeline systems dominate traditional Arabic grammatical analysis: CamelParser2.0 [4], MADAMIRA [9], and Farasa [1] target structured POS+dependency outputs, not prose-form *i’rāb*. We adopt Stanza [10] with a UD→*i’rāb* stub as our published-prior-work baseline, which positions our comparison against the standard non-LLM peer.

Arabic generation models. AraT5v2-base [8], mT5-base [13], AraGPT2-large [3], and AceGPT-13B [6] cover three architectures (T5 enc-dec, GPT-2 decoder-only, Llama-2 decoder-only) and 44 $\times$ parameter range; we train them on the same corpus to isolate architecture and capacity from training data.

Retrieval-augmented prompting. Sonnet RAG and Haiku RAG follow standard retrieval-

augmented prompting [7] with Jaccard retrieval over a 1,060-example pool. Our k -sweep (§3.4) shows the choice of k does no meaningful work on Gazelle.

i’rāb evaluation. **Gazelle** [5] provides the MSA evaluation set; we generate open-text rather than the original closed multiple-choice format. **MASAQ** [12] provides 131K morphological / 123K syntactic Quranic entries with a 72-tag scheme; we use a 5,007-word subset as the cross-register surface and do *not* train on it.

Specialist-generalist routing. MT system combination [11] and similar ensembles motivate our Mix A. To our knowledge, specialist-generalist routing applied to Arabic morphological generation has not been previously evaluated; our Mix A is a direct test of the hypothesis that case/role decompose separately from marker phrasing.

5 Conclusion

We asked two questions about per-word Arabic *i’rāb* generation: does open-weight parameter scale matter when training data is held fixed, and does a system that scores well on MSA news still score well on a different register. The answers are: *not at this scale*, and *no, especially the strongest one*.

Across 296M, 792M, and 13B Arabic-pretrained models trained on the same Haiku-distilled corpus, every pairwise Gazelle Δ has McNemar $p=1.000$. The teacher’s $\sim 67\%$ case ceiling appears to be the binding constraint, not the student’s capacity. The conclusion is conditional on this teacher and on the $n=134$ benchmark; we are explicit about both. Cross-register transfer reveals a sharper picture: every LLM-trained system loses paired-significant role-F1 on Quranic, and Sonnet RAG — our strongest Gazelle system at 32.1% fully-correct — drops by +61.7 pp $\star$, retaining only 19% of its MSA performance. A confound experiment removing the MSA retrieval pool widens the drop, not narrows it. We read this as evidence that the LLMs’ parametric prior is itself MSA-specific, and that within the resolution this benchmark allows, cross-register and construction coverage matter more than parametric capacity.

Future work. The cleanest test of whether the capacity null is teacher-bound is a Sonnet-distilled re-training: an inexpensive single-batch re-distillation that lifts the teacher case ceiling from 67.2% to 73.9% would either keep the open-weight gap null (strong claim) or break it (which would tell us the present null is teacher-bound, not capacity-insensitive). Two further directions follow from §6: expanding Gazelle to a manually-validated 200-sentence MSA benchmark (cuts the smallest-detectable- Δ floor from 7pp to $\sim$

3pp and would let several of our currently-non-significant deltas resolve), and integrating CamelTB / Extended Quranic Treebank into a cross-register training corpus to test whether the cross-register effect is solvable or structural. [2] shows that small fine-tuned Arabic models can be state-of-the-art on diacritization with curated data; the structural connection between diacritization and *i'rāb* (the case marker determines the final vowel; our self-consistency analysis suggests grammatical reasoning constrains diacritization in $\sim 96\%$ of cases) makes a Sadeed-style multi-task formulation a natural extension we leave for future work.

6 Limitations

We restrict our claims as follows. (1) **Sample size.** Gazelle comparisons rest on $n=134$; the smallest detectable Δ at $\alpha=0.05$ power 0.80 is ± 7 pp on a binary metric. AceGPT-13B’s MASAQ row is partial ($n=1,075$, 21%) due to a 4 h SLURM timeout. (2) **Construction coverage.** Verbal (15) and prepositional (8) sentences dominate Gazelle; EXCEPTION and KANA SISTERS appear in 2 sentences each. (3) **Extractor floors.** Marker-EM has a $\sim 10\%$ floor from $\langle \text{NO_MARKER} \rangle$ tokens; role-F1 has the 60% sensitivity floor (§3.4). (4) **Teacher-bound corpus.** All four open-weight FT models inherit Haiku’s case-67% ceiling. The capacity null (§3.2a) holds *given this teacher*. A Sonnet-distilled or human-annotated corpus might reveal capacity differences we cannot detect here. (5) **MASAQ templater asymmetry.** MASAQ gold prose comes from a per-word morphological-tag templater; predictions stay free-form. The subset definition mitigates this gap, but cross-register *fully* comparisons remain weakened. (6) **Closed-source baselines.** Anthropic does not version-pin our calls beyond the `claude-sonnet-4-5` / `claude-haiku-4-5` aliases. Future replications should pin the snapshot. The open-weight FT models reproduce fully from the saved adapters and configs. (7) **Compute.** The 4 h-per-job SLURM limit on the Bocconi MIG 3g.40gb partition forced AceGPT-13B to one epoch and a partial MASAQ run. An A100-80GB schedule would lift both restrictions.

References

- [1] A. Abdelali, K. Darwish, N. Durrani, H. Mubarak. Farasa: A fast and furious segmenter for Arabic.

NAACL Demos, 2016.

- [2] J. Aldallal et al. Sadeed: Small fine-tuned Arabic language models for diacritization. 2025.
- [3] W. Antoun, F. Baly, H. Hajj. AraGPT2: Pre-trained transformer for Arabic language generation. *WANLP*, 2021.
- [4] A. Elshabrawy, N. Habash, et al. CamelParser2.0. *ArabicNLP*, 2023.
- [5] M. Hijjawi et al. Gazelle: An Arabic LLM benchmark. 2024.
- [6] H. Huang et al. AceGPT, localizing large language models in Arabic. *NAACL*, 2024.
- [7] P. Lewis et al. Retrieval-augmented generation for knowledge-intensive NLP tasks. *NeurIPS*, 2020.
- [8] E. M. B. Nagoudi, A. Elmadany, M. Abdul-Mageed. AraT5: Text-to-text transformers for Arabic language generation. *ACL*, 2023.
- [9] A. Pasha et al. MADAMIRA: A fast, comprehensive tool for morphological analysis and disambiguation of Arabic. *LREC*, 2014.
- [10] P. Qi, Y. Zhang, Y. Zhang, J. Bolton, C. D. Manning. Stanza: A Python natural language processing toolkit for many human languages. *ACL Demos*, 2020.
- [11] A.-V. I. Rosti et al. Combining outputs from multiple machine translation systems. *NAACL*, 2007.
- [12] M. Sawalha et al. MASAQ: morphological and syntactic Arabic Quran corpus. 2025.
- [13] L. Xue et al. mT5: A massively multilingual pre-trained text-to-text transformer. *NAACL*, 2021.

A Per-word decoder (negative-result baseline)

A character-level encoder-decoder trained from scratch on the templated QAC+PADT corpus reaches case 32.8% [25.4, 41.0] and role-F1 3.8% [1.5, 8.2] on Gazelle. We report it for completeness only; it does not appear in the main comparison.